

CS460: Intro to Database Systems

*Database System Architectures*

Instructor: Manos Athanassoulis

<http://cs-people.bu.edu/mathan/classes/CS460>

# Today



logistics, goals, admin

when you see this, I  
want you to speak up!  
[and you can always  
interrupt me]

database systems architectures

project details

no smartphones



no laptop



# Course Scope

A detailed look “under the hood” of a DBMS

why?

applications writers, data scientists

database researchers, db admins

they all understand the internals

there is a huge need for database experts

data-intensive applications

big data workflows

# Course Scope: Practical Side

use



benchmark



understand



database systems!

More details when discussing the project!

# Readings

## “Cowbook”

by Ramakrishnan & Gehrke

## Additional Readings

[Architecture of a Database System](#), by J. Hellerstein, M. Stonebraker and J. Hamilton

[The Design and Implementation of Modern Column-store Database Systems](#), by D. Abadi, P. Boncz, S. Harizopoulos, S. Idreos, S. Madden

[Modern B-Tree Techniques](#), by Goetz Graefe, *Foundations and Trends in Databases*, 2011

+research papers



# Guest Lectures

We will have a couple guest lectures

Make sure to attend!

Will be notified ahead of time.



# Evaluation

Class Participation: 5%

## In-class discussion

## Collaborative Notes

3-4 students take notes on shared gdoc

2 days after the class anybody can augment it

<https://tinyurl.com/CS460-f19-notes>

[top part of website as well]

*Enroll right after class!*

# Evaluation

Class Participation: 5%  
Written Assignments: 20%

## **Throughout the semester**

[tentatively] on:

ER model / Relational Model / Relational Algebra

SQL / Normalization

Storage / Disk / Indexing

Transactions / Recovery



# Evaluation

Class Participation: 5%

Written Assignments: 20%

Programming Assignments: 30%

**Three assignments throughout semester**

[more details later today]

# Evaluation

Class Participation: 5%

Written Assignments: 20%

Programming Assignments: 30%

Midterm 1: 20%

Midterm 2: 25%

**both exams during the semester**

# Evaluation

Class Participation: 5%

Written Assignments: 20%

Programming Assignments: 30%

Midterm 1: 20%

Midterm 2: 25%

**SQL Hands-On Test (bonus): 5%**

*Yes! you will use your laptop in class (this once)*


# Office Hours

Manos (before class)

M/W MCS 106 3-4:15pm

TA (will announce in Piazza soon)


# Database Systems



why having a declarative box is useful?



# Database Systems



**application** and **backend** development are independent

collection of algorithms & data structures

multiple ways to do the same thing

**optimization:** dynamically decide which to use

how?



collection of algorithms & data structures

multiple ways to do the same thing

**optimization:** dynamically decide which to use

how? understand & model alternatives



# data management goals



Application

DBMS

DATA



# data management goals



Application



monetary cost




performance

DBMS


DATA



energy



hardware




*“three things are important  
in the database world:  
**performance, performance,  
and performance**”*


Bruce Lindsay, IBM Research

ACM SIGMOD Edgar F. Codd Innovations award 2012

# but



# but



but

new hardware in the last 20 years


multi-core processors

multi-level cache memories

flash drives

SIMD instructions


...



# CS460

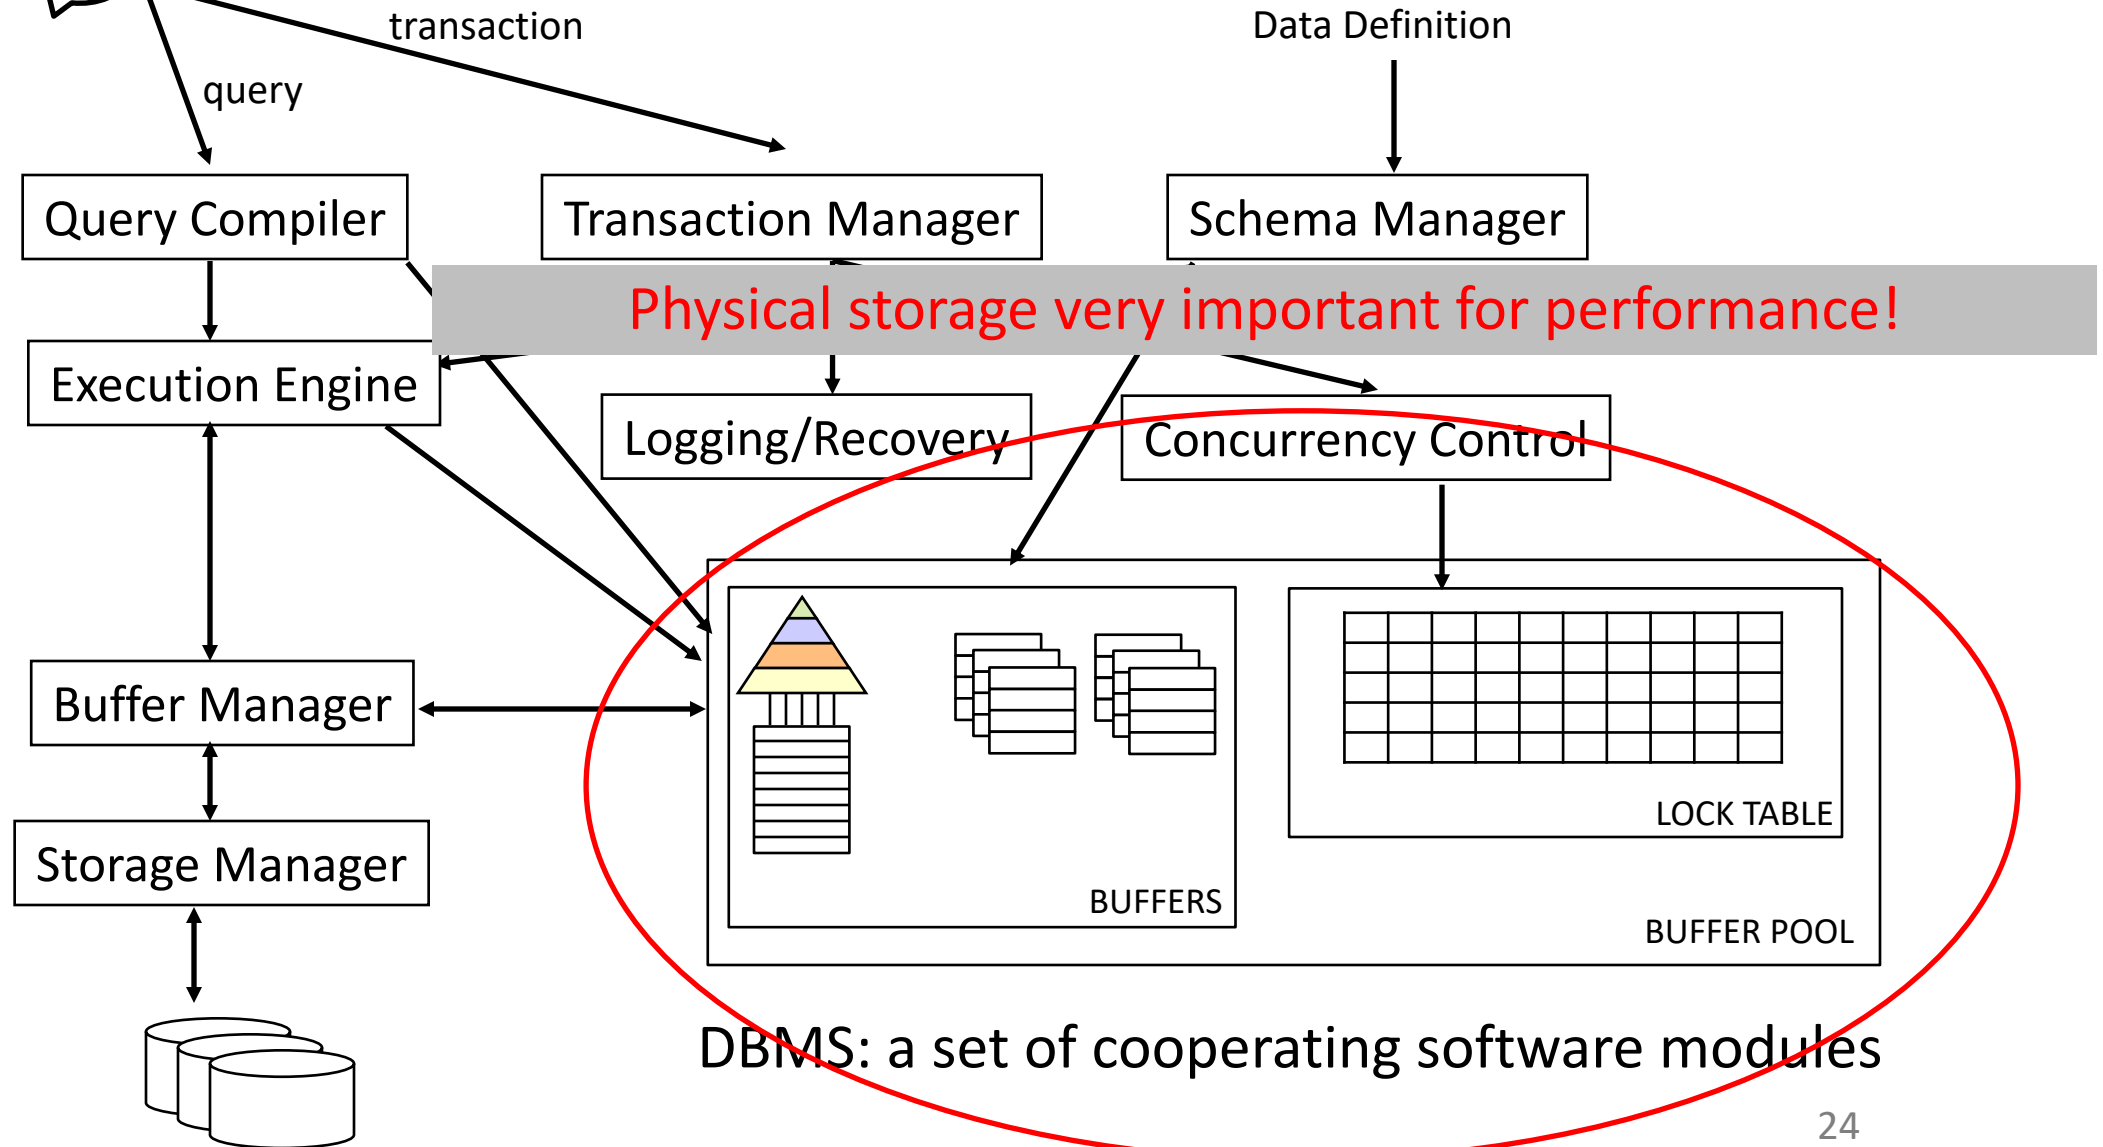
What is inside?

How it works?



performance on  
a declarative box

# Components of a "classic" DBMS





# Some questions for today

how can we physically store our (relational) data?

how to efficiently access the data?

does that affect the way we ***ask*** queries?

does that affect the way we ***evaluate*** queries?

does that affect the way we apply ***updates***?

# how to physically store data?

what is a relation?




a table with rows & columns!

how to physically store it?



# how to physically store data?

one row at a time



# how to efficiently access data?




how to retrieve rows:

if I am interested in the average GPA of all students?

if I am interested in the GPA of student A?

# how to efficiently access data?

Scan the whole table



if I am interested in most of the data

# how to efficiently access data?




how to retrieve rows:

if I am interested in the average GPA of all students?

if I am interested in the GPA of student A?

# how to efficiently access data?

Ask an oracle to tell  
me where is my data




if I am interested in a single row


# how to efficiently access data?

what is an oracle or index?

a data structure that given a value (e.g., student id)

returns location (e.g., row id or a pointer)

with less than  $O(n)$  cost ideally  $O(1)$ !



e.g., B Tree, bitmap, hash index



# how to efficiently access data?

## Scan vs. Index

How to choose?  
Model!

What are the parameters?

data size

index traversal cost

access cost (random vs. sequential)

result set size (“selectivity”)

# how to efficiently access data?

## **Scan vs. Index**

Scan: many rows


Index: few rows

# how to physically store data?


is there another way?



one row at a time




columns first




# how to efficiently access data?

rows first



columns first



if I want to read an entire single row?

if I want to find the name of the younger student?

if I want to calculate the average GPA?

if I want the average GPA of all students with CS Major?

# how to efficiently access data?


## **Rows vs. Columns**

Rows: many attributes+few rows

Columns: few attributes+lots of rows

# does that affect the way we *ask* queries?

I want “blah”



there you go



No!



a declarative box

does that affect the way we ***evaluate*** queries?

Query Engine is different



row-oriented systems ("row-stores")  
move around rows

column-oriented systems ("column-stores")  
move around columns

does that affect the way we ***evaluate*** queries?

ID   Name   DOB   Tel   email   GPA
-------------------------------------


easy mapping from SQL to evaluation strategy

few basic operators: select, project, join, aggregate

simple logic for “query plan”



does that affect the way we ***evaluate*** queries?



simpler basic operators

complicated query logic (more operators to connect)

does that affect the way we apply *updates*?



ID   Name   DOB   Tel   email   GPA	ID	Name	DOB	Tel	email	GPA
ID   Name   DOB   Tel   email   GPA						
ID   Name   DOB   Tel   email   GPA						
ID   Name   DOB   Tel   email   GPA						
ID   Name   DOB   Tel   email   GPA						
ID   Name   DOB   Tel   email   GPA						


how to insert a new row?

how to delete a row?

how to change the GPA of a student?

how to update the email format of all students?

# DBMS timeline



# Row-Stores vs. Column-Stores

physical data layout

simple query plan vs. simple operators

“transactions” vs. “analytics”

# Other Architectures?

## Key-Value Stores (NoSQL)

no transactions

data model: **keys & values**

row: a key and an *arbitrarily complex* value

## Graph Stores

natural representation of graph links

data model: **nodes & relationships**

also maybe: **weights, labels, properties**

# Programming Assignment 1

design, implement, document a database application  
for data, recommendations, reviews for restaurants  
based on real Yelp data

- (1) download & clean
- (2) augment the schema to support additional functionality
  - (3) build an API to the database
  - (4) build a web app that supports:
    - (i) inserting new data, (ii) analysis queries, (iii) browsing

# More Programming Assignments

rows vs. columns (compare the two main paradigms)

query optimization (understand the performance of a query)

key-value systems (deploy and use a KV-system)

# Piazza

Announcements & Discussions in Piazza

<https://piazza.com/bu/fall2019/cs460>





# Remember & Next Time

database systems: performance (energy, HW)

physical storage (row-oriented vs. col-oriented)  
affects query engine/big design space

PA1: build a database application

More programming assignments on

(i) query optimization, (ii) row-stores vs. col-stores, (ii) key-value systems

**Next: Modeling Data**